

Firewall Exploration Lab

学号: 57118225 姓名: 宋雨帆

Task1.A

(1) 将 kernel_module 移动到桌面，并编译

```
[07/24/21] seed@VM:~/.../kernel_module$ make
make -C /lib/modules/5.4.0-54-generic/build M=/home/seed/Desktop/kernel_module modules
make[1]: Entering directory '/usr/src/linux-headers-5.4.0-54-generic'
  CC [M] /home/seed/Desktop/kernel_module/hello.o
  Building modules, stage 2.
  MODPOST 1 modules
WARNING: modpost: missing MODULE_LICENSE() in /home/seed/Desktop/kernel_module/hello.o
see include/linux/module.h for more information
  CC [M] /home/seed/Desktop/kernel_module/hello.mod.o
  LD [M] /home/seed/Desktop/kernel_module/hello.ko
make[1]: Leaving directory '/usr/src/linux-headers-5.4.0-54-generic'
[07/24/21] seed@VM:~/.../kernel_module$
```

(2) 输入指令添加模块

```
[07/24/21] seed@VM:~/.../kernel_module$ sudo insmod hello.ko
[07/24/21] seed@VM:~/.../kernel_module$ lsmod | grep hello
hello                16384  0
[07/24/21] seed@VM:~/.../kernel_module$ sudo rmmod hello
[07/24/21] seed@VM:~/.../kernel_module$ dmesg
```

(3) 安装成功

```
[ 588.308721] Disabling lock debugging due to kernel taint
[ 588.308766] hello: module verification failed: signature and/or required key missing - tainting kernel
[ 588.309952] Hello World!
[ 612.120475] Bye-bye World!.
```

Task1.B.1

(1) 攻击前使用 dig 指令进行查看

```
[07/24/21]seed@VM:~/.../kernel_module$ dig @8.8.8.8 www.example.com

; <<>> DiG 9.16.1-Ubuntu <<>> @8.8.8.8 www.example.com
; (1 server found)
;; global options: +cmd
;; Got answer:
;; ->HEADER<- opcode: QUERY, status: NOERROR, id: 33833
;; flags: qr rd ra ad; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 1

;; OPT PSEUDOSECTION:
; EDNS: version: 0, flags;; udp: 512
;; QUESTION SECTION:
;www.example.com.                IN      A

;; ANSWER SECTION:
www.example.com.                20451   IN      A      93.184.216.34

;; Query time: 44 msec
;; SERVER: 8.8.8.8#53(8.8.8.8)
;; WHEN: Sat Jul 24 15:05:57 EDT 2021
;; MSG SIZE rcvd: 60
```

(2) 编译 packet_filter 内核

```
[07/24/21]seed@VM:~/.../packet_filter$ make
make -C /lib/modules/5.4.0-54-generic/build M=/home/seed/Desktop/packet_filter modules
make[1]: Entering directory '/usr/src/linux-headers-5.4.0-54-generic'
  CC [M] /home/seed/Desktop/packet_filter/seedFilter.o
  Building modules, stage 2.
  MODPOST 1 modules
  CC [M] /home/seed/Desktop/packet_filter/seedFilter.mod.o
  LD [M] /home/seed/Desktop/packet_filter/seedFilter.ko
make[1]: Leaving directory '/usr/src/linux-headers-5.4.0-54-generic'
```

(3) 输入指令加载模块

```
[07/24/21]seed@VM:~/.../packet_filter$ sudo insmod seedFilter.ko
[07/24/21]seed@VM:~/.../packet_filter$ lsmod | grep seedFilter
seedFilter                16384  0
```

(4) 重新输入 dig 指令

```
[07/24/21]seed@VM:~/.../kernel_module$ dig @8.8.8.8 www.example.com

; <<>> DiG 9.16.1-Ubuntu <<>> @8.8.8.8 www.example.com
; (1 server found)
;; global options: +cmd
;; connection timed out; no servers could be reached
```

可以看到无法连接

Task1.B.2

(1) 修改程序，新增 hook 并将其与 printInfo 关联

```
static struct nf_hook_ops hook1, hook2, hook3, hook4, hook5;
```



```

1 int registerFilter(void) {
2     printk(KERN_INFO "Registering filters.\n")
3
4     hook1.hook = printInfo;
5     hook1.hooknum = NF_INET_PRE_ROUTING;
6     hook1.pf = PF_INET;
7     hook1.priority = NF_IP_PRI_FIRST;
8     nf_register_net_hook(&init_net, &hook1);
9
10    hook2.hook = printInfo;
11    hook2.hooknum = NF_INET_LOCAL_IN;
12    hook2.pf = PF_INET;
13    hook2.priority = NF_IP_PRI_FIRST;
14    nf_register_net_hook(&init_net, &hook2);
15
16    hook3.hook = printInfo;
17    hook3.hooknum = NF_INET_FORWARD;
18    hook3.pf = PF_INET;
19    hook3.priority = NF_IP_PRI_FIRST;
20    nf_register_net_hook(&init_net, &hook3);
21
22    hook4.hook = printInfo;
23    hook4.hooknum = NF_INET_LOCAL_OUT;
24    hook4.pf = PF_INET;
25    hook4.priority = NF_IP_PRI_FIRST;
26    nf_register_net_hook(&init_net, &hook4);
27
28    hook5.hook = printInfo;
29    hook5.hooknum = NF_INET_POST_ROUTING;
30    hook5.pf = PF_INET;
31    hook5.priority = NF_IP_PRI_FIRST;
32    nf_register_net_hook(&init_net, &hook5);
33
34    void removeFilter(void) {
35        printk(KERN_INFO "The filters are being removed.\n");
36        nf_unregister_net_hook(&init_net, &hook1);
37        nf_unregister_net_hook(&init_net, &hook2);
38        nf_unregister_net_hook(&init_net, &hook3);
39        nf_unregister_net_hook(&init_net, &hook4);
40        nf_unregister_net_hook(&init_net, &hook5);
41    }
}

```

(2) 重新编译和加载

```
[07/24/21]seed@VM:~/.../packet_filter$ make
make -C /lib/modules/5.4.0-54-generic/build M=/home/seed/Desktop/packet_filter modules
make[1]: Entering directory '/usr/src/linux-headers-5.4.0-54-generic'
Building modules, stage 2.
MODPOST 1 modules
LD [M] /home/seed/Desktop/packet_filter/seedFilter.ko
make[1]: Leaving directory '/usr/src/linux-headers-5.4.0-54-generic'
[07/24/21]seed@VM:~/.../packet_filter$ sudo insmod seedFilter.ko
insmod: ERROR: could not insert module seedFilter.ko: File exists
[07/24/21]seed@VM:~/.../packet_filter$ sudo rmmod seedFilter.ko
[07/24/21]seed@VM:~/.../packet_filter$ sudo insmod seedFilter.ko
[07/24/21]seed@VM:~/.../packet_filter$ lsmod | grep seed
seedFilter                16384  0
```

(3) ping 内网，输入 dmesg 查看日志

```
[07/24/21]seed@VM:~/.../kernel_module$ ping 192.168.214.2
PING 192.168.214.2 (192.168.214.2) 56(84) bytes of data.
64 bytes from 192.168.214.2: icmp_seq=1 ttl=128 time=0.295 ms
64 bytes from 192.168.214.2: icmp_seq=2 ttl=128 time=0.632 ms
64 bytes from 192.168.214.2: icmp_seq=3 ttl=128 time=0.520 ms
64 bytes from 192.168.214.2: icmp_seq=4 ttl=128 time=0.751 ms
```

可以看到，对于本地产生向外部网络发出的数据包，会依次经过 NF_INET_LOCAL_OUT 和 NF_INET_POST_ROUTING 两个 hook 点的处理。推测前者是本地产生包时触发，后者是向外发送时触发。

```
*** LOCAL_OUT
192.168.214.128 --> 192.168.214.2 (ICMP)
*** POST_ROUTING
192.168.214.128 --> 192.168.214.2 (ICMP)
```

发往本地的数据包，会依次经过 NF_INET_PRE_ROUTING 和 NF_INET_LOCAL_IN 两个 hook 点的处理。推测前者是判断是发往外部还是本地接收时触发，后者是发往本地时触发。

```
*** PRE_ROUTING
192.168.214.2 --> 192.168.214.128 (ICMP)
*** LOCAL_IN
192.168.214.2 --> 192.168.214.128 (ICMP)
```

在日志中没有出现 NF_INET_FORWARD 点的活动

因此得到结论：

NF_INET_LOCAL_OUT	触发条件：由本地产生的包
NF_INET_POST_ROUTING	触发条件：向外部网络发出的包
NF_INET_PRE_ROUTING	触发条件：本地接收到的包（不包括本地产生的），用于判断是否向外转发
NF_INET_LOCAL_IN	触发条件：发往本地而不需要转发的包
NF_INET_FORWARD	触发条件：需要向外转发的数据包

Task1.B.3

(1) 增加两个 hook 函数，分别用于拦截 icmp 报文和 telnet 报文

```
unsigned int blockicmp(void*priv, struct sk_buff *skb, const struct nf_hook_state *state)
{
    struct iphdr *iph;

    iph = ip_hdr(skb);

    if(iph->protocol == IPPROTO_ICMP)
    {
        printk(KERN_WARNING "*** Dropping %pI4 (ICMP)\n", &(iph->daddr));
        return NF_DROP;
    }
    return NF_ACCEPT;
}

unsigned int blocktelnet(void*priv, struct sk_buff *skb, const struct nf_hook_state *state)
{
    struct iphdr *iph;
    struct tcphdr *tcph;

    u16 port = 23;
    iph = ip_hdr(skb);
    if(iph->protocol == IPPROTO_TCP)
    {
        tcph = tcp_hdr(skb);
        if(ntohs(tcph->dest)==port)
        {
            printk(KERN_WARNING "*** Dropping %pI4 (TCP), port %d\n", &(iph->daddr), port);
            return NF_DROP;
        }
    }
    return NF_ACCEPT;
}
```

其中可以看到，这里没有判断目标 ip 是否为主机，这是因为利用 task1.B.2 中的知识，所有触发 NF_INET_LOCAL_IN 的都是发往本地的包，因此无需再比较 ip，只需要将函数注册为 NF_INET_LOCAL_IN 点即可。

```
int registerFilter(void) {
    printk(KERN_INFO "Registering filters.\n");

    hook1.hook = blockicmp;
    hook1.hooknum = NF_INET_LOCAL_IN;
    hook1.pf = PF_INET;
    hook1.priority = NF_IP_PRI_FIRST;
    nf_register_net_hook(&init_net, &hook1);

    hook2.hook = blocktelnet;
    hook2.hooknum = NF_INET_LOCAL_IN;
    hook2.pf = PF_INET;
    hook2.priority = NF_IP_PRI_FIRST;
    nf_register_net_hook(&init_net, &hook2);

    return 0;
}
```

(2) 重新编译并加载模块


```
[07/24/21]seed@VM:~/.../packet_filter$ sudo rmmod seedFilter.ko
[07/24/21]seed@VM:~/.../packet_filter$ sudo insmod seedFilter.ko
[07/24/21]seed@VM:~/.../packet_filter$ lsmod | grep seed
seedFilter                16384  0
```

(3) 登录 hostA 主机，并 ping 10.9.0.1（虚拟机），可以看到没有 ping 通

```
root@a2ea5eaaa347:/# ping 10.9.0.1
PING 10.9.0.1 (10.9.0.1) 56(84) bytes of data.
```

输入 dmesg 查看

```
[10788.259120] *** Dropping 10.9.0.1 (ICMP)
[10789.273805] *** Dropping 10.9.0.1 (ICMP)
[10790.300212] *** Dropping 10.9.0.1 (ICMP)
[10791.322934] *** Dropping 10.9.0.1 (ICMP)
[10792.346717] *** Dropping 10.9.0.1 (ICMP)
[10793.369339] *** Dropping 10.9.0.1 (ICMP)
[10794.395158] *** Dropping 10.9.0.1 (ICMP)
[10795.416869] *** Dropping 10.9.0.1 (ICMP)
[10796.443848] *** Dropping 10.9.0.1 (ICMP)
[10797.466153] *** Dropping 10.9.0.1 (ICMP)
```

可以看到，icmp 报文 都被丢弃了，说明拦截成功

(4) 使用 telnet 10.9.0.1 尝试登录，可以看到登录失败

```
root@a2ea5eaaa347:/# telnet 10.9.0.1
Trying 10.9.0.1...
```

输入 dmesg 命令查看

```
[10814.489442] *** Dropping 10.9.0.1 (TCP), port 23
[10814.940865] *** Dropping 10.9.0.1 (TCP), port 23
[10815.961314] *** Dropping 10.9.0.1 (TCP), port 23
[10817.978420] *** Dropping 10.9.0.1 (TCP), port 23
[10822.168420] *** Dropping 10.9.0.1 (TCP), port 23
[10830.362918] *** Dropping 10.9.0.1 (TCP), port 23
[10846.489516] *** Dropping 10.9.0.1 (TCP), port 23
[10880.027292] *** Dropping 10.9.0.1 (TCP), port 23
```

可以看到，丢弃了很多 tcp 报文，从而阻止了登录，说明拦截成功。

Task2.A

(1) 未进行操作前，在主机上 ping 路由器，成功

```
root@a2ea5eaaa347:/# ping 10.9.0.11
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.
64 bytes from 10.9.0.11: icmp_seq=1 ttl=64 time=0.219 ms
64 bytes from 10.9.0.11: icmp_seq=2 ttl=64 time=0.124 ms
64 bytes from 10.9.0.11: icmp_seq=3 ttl=64 time=0.152 ms
64 bytes from 10.9.0.11: icmp_seq=4 ttl=64 time=0.219 ms
64 bytes from 10.9.0.11: icmp_seq=5 ttl=64 time=0.106 ms
64 bytes from 10.9.0.11: icmp_seq=6 ttl=64 time=0.110 ms
64 bytes from 10.9.0.11: icmp_seq=7 ttl=64 time=0.166 ms
```

(2) 输入规则

```
root@6215949f23c8:/# iptables -A OUTPUT -p icmp --icmp-type echo-reply -j ACCEPT
root@6215949f23c8:/# iptables -A INPUT -p icmp --icmp-type echo-request -j ACCEPT
root@6215949f23c8:/# iptables -t filter -L -n
Chain INPUT (policy ACCEPT)
target     prot opt source                destination            icmptype 8
ACCEPT     icmp -- 0.0.0.0/0              0.0.0.0/0

Chain FORWARD (policy ACCEPT)
target     prot opt source                destination

Chain OUTPUT (policy ACCEPT)
target     prot opt source                destination            icmptype 0
ACCEPT     icmp -- 0.0.0.0/0              0.0.0.0/0
root@6215949f23c8:/# iptables -P OUTPUT DROP
root@6215949f23c8:/# iptables -P INPUT DROP
```

(3) 测试

使用 ping 指令，成功

```
root@a2ea5eaaa347:/# ping 10.9.0.11
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.
64 bytes from 10.9.0.11: icmp_seq=1 ttl=64 time=0.130 ms
64 bytes from 10.9.0.11: icmp_seq=2 ttl=64 time=0.052 ms
64 bytes from 10.9.0.11: icmp_seq=3 ttl=64 time=0.147 ms
64 bytes from 10.9.0.11: icmp_seq=4 ttl=64 time=0.131 ms
64 bytes from 10.9.0.11: icmp_seq=5 ttl=64 time=0.152 ms
64 bytes from 10.9.0.11: icmp_seq=6 ttl=64 time=0.150 ms
```

使用 telnet 指令，失败

```
root@a2ea5eaaa347:/# telnet 10.9.0.11
Trying 10.9.0.11...
telnet: Unable to connect to remote host: Connection timed out
说明规则编写正确。
```

Task2.B

(1) 在路由器上输入 ifconfig 查看端口对应 ip


```

root@6215949f23c8:/# ifconfig
eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.9.0.11 netmask 255.255.255.0 broadcast 10.9.0.255
    ether 02:42:0a:09:00:0b txqueuelen 0 (Ethernet)
    RX packets 81 bytes 8119 (8.1 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 18 bytes 1484 (1.4 KB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.60.11 netmask 255.255.255.0 broadcast 192.168.60.255
    ether 02:42:c0:a8:3c:0b txqueuelen 0 (Ethernet)
    RX packets 55 bytes 6019 (6.0 KB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    loop txqueuelen 1000 (Local Loopback)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

```

可以看到 10.9.0.11 对应的端口(eth0)是面向外网的, 192.168.60.11(eth1)对应的端口是面向内网的。

(2) 根据要求指定规则如下:

①OUTPUT INPUT FORWARD 丢弃

```
iptables -P OUTPUT DROP
```

```
iptables -P INPUT DROP
```

```
iptables -P FORWARD DROP
```

(不允许内外流量交互)

②对于 FORWARD 只有 icmp 请求报文由内部端口 (eth1) 进入, 外部端口 (eth0), 才接收:

```
iptables -A FORWARD -p icmp --icmp-type echo-request -i eth1 -j ACCEPT
```

```
iptables -A FORWARD -p icmp --icmp-type echo-request -o eth0 -j ACCEPT
```

对于 FORWARD 只有 icmp 应答报文由外部端口 (eth0) 进入, 外部端口 (eth1), 才接收:

```
iptables -A FORWARD -p icmp --icmp-type echo-reply -i eth0 -j ACCEPT
```

```
iptables -A FORWARD -p icmp --icmp-type echo-reply -o eth1 -j ACCEPT
```

(从而保证外部不能 ping 内部, 内部能 ping 外部)

③对于 input 和 output, 允许输入的 icmp 请求和应答报文, 保证外部主机能够 ping 路由器

```
iptables -A INPUT -p icmp --icmp-type echo-request -j ACCEPT
```

```
iptables -A OUTPUT -p icmp --icmp-type echo-reply -j ACCEPT
```

```

root@6215949f23c8:/# iptables -A INPUT -p icmp --icmp-type echo-request -j ACCEPT
root@6215949f23c8:/# iptables -A OUTPUT -p icmp --icmp-type echo-reply -j ACCEPT
root@6215949f23c8:/# iptables -A FORWARD -p icmp --icmp-type echo-request -i eth1 -j ACCEPT
root@6215949f23c8:/# iptables -A FORWARD -p icmp --icmp-type echo-request -o eth0 -j ACCEPT
root@6215949f23c8:/# iptables -A FORWARD -p icmp --icmp-type echo-reply -i eth0 -j ACCEPT
root@6215949f23c8:/# iptables -A FORWARD -p icmp --icmp-type echo-reply -o eth1 -j ACCEPT
root@6215949f23c8:/# iptables -P OUTPUT DROP
root@6215949f23c8:/# iptables -P INPUT DROP
root@6215949f23c8:/# iptables -P FORWARD DROP

```


(3) 测试

①外网 ping 内网

```
root@a2ea5eaaa347:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
```



不能实现 ping 操作，符合要求。

②内网 ping 外网

```
[07/24/21]seed@VM:~/Desktop$ docksh 64
root@64c563db5186:/# ping 10.9.0.5
PING 10.9.0.5 (10.9.0.5) 56(84) bytes of data.
64 bytes from 10.9.0.5: icmp_seq=1 ttl=63 time=0.328 ms
64 bytes from 10.9.0.5: icmp_seq=2 ttl=63 time=0.193 ms
64 bytes from 10.9.0.5: icmp_seq=3 ttl=63 time=0.076 ms
64 bytes from 10.9.0.5: icmp_seq=4 ttl=63 time=0.186 ms
```

能够实现 ping 操作，符合要求。

③外网 ping 路由

```
root@a2ea5eaaa347:/# ping 10.9.0.11
PING 10.9.0.11 (10.9.0.11) 56(84) bytes of data.
64 bytes from 10.9.0.11: icmp_seq=1 ttl=64 time=0.166 ms
64 bytes from 10.9.0.11: icmp_seq=2 ttl=64 time=0.131 ms
64 bytes from 10.9.0.11: icmp_seq=3 ttl=64 time=0.148 ms
64 bytes from 10.9.0.11: icmp_seq=4 ttl=64 time=0.152 ms
```

能够实现 ping 操作，符合要求。

④内外其它流量的交互

外网 telnet 内网:

```
root@a2ea5eaaa347:/# telnet 192.168.60.5
Trying 192.168.60.5...
```



无法登录，符合要求。

内网 telnet 外网:

```
root@64c563db5186:/# telnet 10.9.0.5
Trying 10.9.0.5...
```

无法登录，符合要求。

Task2.C

(1) 设计策略如下

```
root@6215949f23c8:/# iptables -A FORWARD -p tcp -d 192.168.60.5 --dport 23 -j ACCEPT
root@6215949f23c8:/# iptables -A FORWARD -p tcp -s 192.168.60.5 --sport 23 -j ACCEPT
root@6215949f23c8:/# iptables -P FORWARD DROP
```

只允许 192.168.60.5 主机的 23 端口的流量进行转发, 保证其能被外部和内部主机登录, 而外部主机到其它内部主机则不行; 本地主机的相互访问不需要转发, 因此 FORWARD 对其它报文的丢弃不会影响内部主机间的登录。

①外部 telnet 192.168.60.5

```
root@a2ea5eaaa347:/# telnet 192.168.60.5
Trying 192.168.60.5...
Connected to 192.168.60.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
64c563db5186 login: █
```

成功。

②外部 telnet 其它主机

```
root@a2ea5eaaa347:/# telnet 192.168.60.7
Trying 192.168.60.7...
```

无法登录, 符合要求。

③内部 telnet 外部

```
root@64c563db5186:/# telnet 10.9.0.5
Trying 10.9.0.5...
```

无法登录, 符合要求。

④内部 telnet 内部

```
root@64c563db5186:/# telnet 192.168.60.6
Trying 192.168.60.6...
Connected to 192.168.60.6.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
b9ae5085eeaa login:
```

成功。

Task3.A

ICMP

(1) 在 10.9.0.5 上 ping 192.168.60.5 主机

```
root@c5d90be81534:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.323 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.180 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.065 ms
64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.074 ms
```


(2) 查看追踪信息

可以看到一个 icmp。

```
root@6215949f23c8:/# conntrack -L
icmp      1 2 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=53 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=53 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
root@6215949f23c8:/# conntrack -L
conntrack v1.4.5 (conntrack-tools): 0 flow entries have been shown.
```

如果 30s 内再次 ping，则会显示两个 flow。

```
root@6215949f23c8:/# conntrack -L
icmp      1 22 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=54 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=54 mark=0 use=1
icmp      1 29 src=10.9.0.5 dst=192.168.60.5 type=8 code=0 id=55 src=192.168.60.5 dst=10.9.0.5 type=0 code=0 id=55 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 2 flow entries have been shown.
```

UDP

(1) 在主机 10.9.0.5 上输入指令

```
root@a2ea5eaaa347:/# nc -u 192.168.60.5 9090
```

(2) 查看追踪信息

```
root@6215949f23c8:/# conntrack -L
udp       17 9 src=10.9.0.5 dst=192.168.60.5 sport=43197 dport=9090 [UNREPLIED] src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=43197 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
```

时间同样为 30s，且停止登录后会倒计时到 0。

```
root@6215949f23c8:/# conntrack -L
conntrack v1.4.5 (conntrack-tools): 0 flow entries have been shown.
```

同样，30s 内再次连接会出现两条记录。

```
root@6215949f23c8:/# conntrack -L
udp       17 9 src=10.9.0.5 dst=192.168.60.5 sport=40028 dport=9090 [UNREPLIED] src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=40028 mark=0 use=1
udp       17 26 src=10.9.0.5 dst=192.168.60.5 sport=47850 dport=9090 [UNREPLIED] src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=47850 mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 2 flow entries have been shown.
```

TCP

(1) 在当前机和目标主机上输入指令，建立连接

```
root@c5d90be81534:/# nc 192.168.60.5 9090
haha
```

```
root@8c73b014b688:/# nc -l 9090
haha
_
```

(2) 查看追踪

```
root@6215949f23c8:/# conntrack -L
tcp        6 431995 ESTABLISHED src=10.9.0.5 dst=192.168.60.5 sport=55050 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=55050 [ASSURED] mark=0 use=1
conntrack v1.4.5 (conntrack-tools): 1 flow entries have been shown.
```

可以看到，存活时间在 430000s 以上

(3) 结束连接，可以看到记录仍然存在，但只存在 120s，120s 倒计时结束后将消失。

```
|tcp        6 113 TIME WAIT src=10.9.0.5 dst=192.168.60.5 sport=55054 dport=9090 src=192.168.60.5 dst=10.9.0.5 sport=9090 dport=55054 [ASSURED] mark=0 use=1
```

Task3.B

(1) 编写规则如下

```
root@6215949f23c8:/# iptables -A FORWARD -p tcp -m conntrack --ctstate ESTABLISHED,RELATED -j ACCEPT
root@6215949f23c8:/# iptables -A FORWARD -p tcp -m conntrack --ctstate NEW -i eth1 -j ACCEPT
root@6215949f23c8:/# iptables -P FORWARD DROP
```

(2) 测试

内网 telnet 外网，成功

```
root@64c563db5186:/# telnet 10.9.0.5
Trying 10.9.0.5...
Connected to 10.9.0.5.
Escape character is '^]'.
Ubuntu 20.04.1 LTS
a2ea5eaaa347 login: seed
Password:
Welcome to Ubuntu 20.04.1 LTS (GNU/Linux 5.4.0-54-generic x86_64)
```

外网 telnet 内网，无法连接

```
root@a2ea5eaaa347:/# telnet 192.168.60.6
Trying 192.168.60.6...
```

因此规则正确，符合实验要求。

Task4

(1) 在路由器输入规则

```
root@6215949f23c8:/# iptables -A FORWARD -s 10.9.0.5 -m limit --limit 10/minute --limit-burst 5 -j ACCEPT
root@6215949f23c8:/# iptables -A FORWARD -s 10.9.0.5 -j DROP
```

(2) 在 10.9.0.5 主机尝试 ping

```
root@a2ea5eaaa347:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.544 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.192 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.166 ms
64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.290 ms
64 bytes from 192.168.60.5: icmp_seq=5 ttl=63 time=0.150 ms
64 bytes from 192.168.60.5: icmp_seq=7 ttl=63 time=0.184 ms
64 bytes from 192.168.60.5: icmp_seq=13 ttl=63 time=0.178 ms
64 bytes from 192.168.60.5: icmp_seq=19 ttl=63 time=0.187 ms
64 bytes from 192.168.60.5: icmp_seq=25 ttl=63 time=0.167 ms
```

可以看到，前几个报文速度非常快，后面速度较慢，平均 6s 一个，符合要求，说明规则发挥了作用

(3) 去掉第二条规则，再次尝试 ping

```
root@6215949f23c8:/# iptables -F
root@6215949f23c8:/# iptables -A FORWARD -s 10.9.0.5 -m limit --limit 10/minute --limit-burst 5 -j ACCEPT
```



```

root@a2ea5eaaa347:/# ping 192.168.60.5
PING 192.168.60.5 (192.168.60.5) 56(84) bytes of data.
64 bytes from 192.168.60.5: icmp_seq=1 ttl=63 time=0.246 ms
64 bytes from 192.168.60.5: icmp_seq=2 ttl=63 time=0.063 ms
64 bytes from 192.168.60.5: icmp_seq=3 ttl=63 time=0.123 ms
64 bytes from 192.168.60.5: icmp_seq=4 ttl=63 time=0.208 ms
64 bytes from 192.168.60.5: icmp_seq=5 ttl=63 time=0.192 ms
64 bytes from 192.168.60.5: icmp_seq=6 ttl=63 time=0.168 ms
64 bytes from 192.168.60.5: icmp_seq=7 ttl=63 time=0.075 ms
64 bytes from 192.168.60.5: icmp_seq=8 ttl=63 time=0.171 ms
64 bytes from 192.168.60.5: icmp_seq=9 ttl=63 time=0.360 ms
64 bytes from 192.168.60.5: icmp_seq=10 ttl=63 time=0.195 ms
64 bytes from 192.168.60.5: icmp_seq=11 ttl=63 time=0.225 ms

```

可以看到，报文发送速度很快，并没有减慢到每分钟 10 个，说明规则失效了。

这是因为没有第二条规则将报文设置为默认 DROP，所有报文（无论是否满足第一个规则）都会从默认的 ACCEPT 规则通过，从而导致第一条规则也失效。

Task5

轮询：

（1）在路由器上输入规则

```

root@6215949f23c8:/# iptables -F
root@6215949f23c8:/# iptables -t nat -A PREROUTING -p udp --dport 8080 \
> -m statistic --mode nth --every 3 --packet 0 \
> -j DNAT --to-destination 192.168.60.5:8080

```

（2）在 192.168.60.5 上开启 nc -luk 8080 监听，在 10.9.0.5 处连接并输入 hello

```

root@a2ea5eaaa347:/# echo hello | nc -u 10.9.0.11 8080
^C
root@a2ea5eaaa347:/# echo hello | nc -u 10.9.0.11 8080
root@a2ea5eaaa347:/# echo hello | nc -u 10.9.0.11 8080

```

（3）在 192.168.60.5 上查看，可以看到发送者没发送 3 次，接收者再会收到 1 次路由器发的报文，说明路由规则正确。

```

root@64c563db5186:/# nc -luk 8080
hello

```

随机：

（1）指定规则如下

```

root@6215949f23c8:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode random --probability 0.333 -j DNAT --to-destination 192.168.60.5:8080
root@6215949f23c8:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -m statistic --mode random --probability 0.5 -j DNAT --to-destination 192.168.60.6:8080
root@6215949f23c8:/# iptables -t nat -A PREROUTING -p udp --dport 8080 -j DNAT --to-destination 192.168.60.7:8080

```

其中第一个的概率为 0.333，第二个概率为 0.5，当前两个没有命中，则默认命中第三条

（2）在三台主机上都输入 nc -luk 8080 打开监听，在 10.9.0.5 主机上建立连接并不断输入 hello，得到三台主机上的输出结果如下：

```
root@28278760de5d:/# nc -luk 8080
hello
hello
hello
hello
hello
hello
hello
hello
hello
hello
hello
hello
```

```
root@b9ae5085eeaa:/# nc -luk 8080
hello
hello
hello
hello
hello
hello
hello
hello
hello
hello
hello
hello
```

```
root@64c563db5186:/# nc -luk 8080
hello
hello
hello
hello
hello
hello
hello
hello
hello
hello
hello
hello
hello
hello
hello
```

可以看到路由确实实现了分发, 由于样本数量太少, 不能做到每个主机上收到的报文各 1/3, 但总体上符合负载均衡的要求。